

The Data Driven Approach to predict the success of Bank Telemarketing

2023-05-08

Data information

The data is related with direct marketing campaigns of a Portuguese banking institution. The marketing campaigns were based on phone calls. Often more than one contact to the same client was required.

The Business problem I am trying to solve

- The goal of this project is to predict the customer would subscribed bank term deposit or not.
- In order to achieve this objective, first we need to explore and compare which model gives the accurate and better results.
- We will be examining parametric method such as logistic regression, as well as non parametric method such as KNN, to determine the most effective approach.

Load the library

load the dataset

```
Bank=read.csv("~/Desktop/SFSU/math449Project/bank.csv",header=TRUE,sep=";")
```

To see the first five row of the data

```
head(Bank)
```

```
##   age      job marital education default balance housing loan  contact day
## 1  30  unemployed married  primary      no   1787      no   no cellular  19
## 2  33   services married secondary     no   4789     yes  yes cellular  11
## 3  35  management single  tertiary     no   1350     yes   no cellular  16
## 4  30  management married  tertiary     no   1476     yes  yes  unknown   3
## 5  59 blue-collar married secondary     no     0      yes   no  unknown   5
## 6  35  management single  tertiary     no    747      no   no cellular  23
##  month duration campaign pdays previous poutcome y
## 1  oct         79         1    -1         0  unknown no
## 2  may        220         1   339         4  failure no
## 3  apr        185         1   330         1  failure no
## 4  jun        199         4    -1         0  unknown no
## 5  may        226         1    -1         0  unknown no
## 6  feb        141         2   176         3  failure no
```

to see the data type of Bank.csv

```
str(Bank)
```

```
## 'data.frame':   4521 obs. of  17 variables:
## $ age      : int  30 33 35 30 59 35 36 39 41 43 ...
## $ job      : chr  "unemployed" "services" "management" "management" ...
## $ marital  : chr  "married" "married" "single" "married" ...
## $ education: chr  "primary" "secondary" "tertiary" "tertiary" ...
## $ default  : chr  "no" "no" "no" "no" ...
## $ balance  : int  1787 4789 1350 1476 0 747 307 147 221 -88 ...
## $ housing  : chr  "no" "yes" "yes" "yes" ...
## $ loan     : chr  "no" "yes" "no" "yes" ...
## $ contact  : chr  "cellular" "cellular" "cellular" "unknown" ...
## $ day      : int  19 11 16 3 5 23 14 6 14 17 ...
## $ month    : chr  "oct" "may" "apr" "jun" ...
## $ duration : int  79 220 185 199 226 141 341 151 57 313 ...
## $ campaign : int  1 1 1 4 1 2 1 2 2 1 ...
## $ pdays   : int  -1 339 330 -1 -1 176 330 -1 -1 147 ...
## $ previous : int  0 4 1 0 0 3 2 0 0 2 ...
## $ poutcome : chr  "unknown" "failure" "failure" "unknown" ...
## $ y        : chr  "no" "no" "no" "no" ...
```


correlatin between numeric predictors variable

```
BankNumeric=select_if(Bank,is.integer)
corMatrix=cor(BankNumeric)
corrplot(corMatrix,type="upper",method="color",tl.col="black",tl.srt=45)
```

```
## Warning in text.default(pos.xlabel[, 1], pos.xlabel[, 2], newcolnames, srt =  
## t1.srt, : "t1.srt" is not a graphical parameter
```

```
## Warning in text.default(pos.ylabel[, 1], pos.ylabel[, 2], newrownames, col =  
## t1.col, : "t1.srt" is not a graphical parameter
```

```
## Warning in title(title, ...): "t1.srt" is not a graphical parameter
```



The correlation coefficients between all pairs of predictor variables in the model are less than 0.5. Thus, we don't need to worry about multicollinearity in this problem.

count for non-numeric values

```
cols=c("job","marital","education","default","housing","contact","month","poutcome")
for (col in cols){
  counts=table(Bank[,col][!is.numeric(Bank[,col])])
  cat(paste0("Counts for ",col, " column :\n"))
  print(counts)
  cat("\n")
}
```

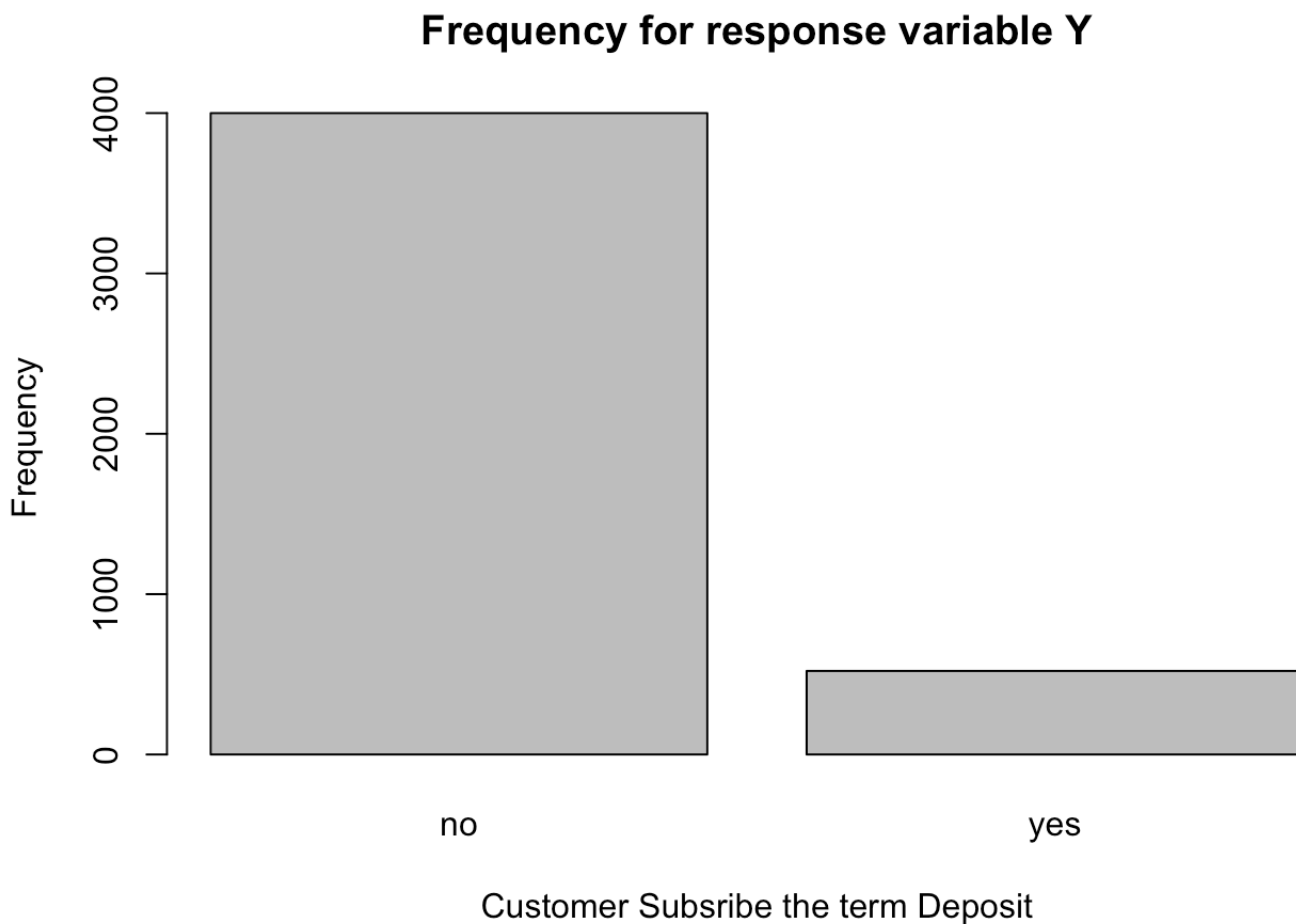
```
## Counts for job column :
##
##      admin.   blue-collar  entrepreneur   housemaid   management
##      478      946      168      112      969
##      retired self-employed   services      student   technician
##      230      183      417      84      768
##      unemployed      unknown
##      128      38
##
## Counts for marital column :
##
## divorced  married   single
##      528      2797      1196
##
## Counts for education column :
##
##      primary secondary  tertiary   unknown
##      678      2306      1350      187
##
## Counts for default column :
##
##      no  yes
## 4445   76
##
## Counts for housing column :
##
##      no  yes
## 1962 2559
##
## Counts for contact column :
##
##      cellular telephone   unknown
##      2896      301      1324
##
## Counts for month column :
##
##      apr  aug  dec  feb  jan  jul  jun  mar  may  nov  oct  sep
##      293  633   20  222  148  706  531  49  1398  389   80   52
##
## Counts for poutcome column :
##
##      failure   other  success  unknown
##      490      197      129      3705
```

Graph the non numeric value counts for each column

```
dev.new( )  
for (col in 1:length(cols)){  
  counts=table(Bank[,cols[col]][!is.numeric(Bank[,cols[col]])])  
  barplot(counts,main=cols[col],xlab="Non Numeric values", ylab="Count")  
}
```

value counts for response variable

```
countsY=table(Bank$y)  
barplot(countsY,main="Frequency for response variable Y",xlab="Customer Subscribe the  
term Deposit",ylab="Frequency")
```



compute the percentage of yes and no

```
percentageYes=countsY[2] / nrow(Bank) * 100
percentageNo=countsY[1] / nrow(Bank) * 100

cat(paste0("Percentage of subscription of term deposit: ",format(percentageYes,nsml
=2), " %"))
```

```
## Percentage of subscription of term deposit: 11.524 %
```

```
cat("\n")
```

```
cat(paste0("Percentage of no subscription of term deposit: ",format(percentageNo,nsma
ll=2), " %"))
```

```
## Percentage of no subscription of term deposit: 88.476 %
```

Based on the graph and compute percentage, the response variable is unbalanced, with more “No” response than “Yes” which could cause biased in our prediction. This means it may accurately predict the majority class which is “No”, but fail to accurately predict the minority class.

I will try to make it balance by using the oversampling method

```
Bank=ovun.sample(y~.,data=Bank,method="both",N=nrow(Bank),seed=123)$data
table(Bank$y)
```

```
##
##   no   yes
## 2279 2242
```

Now, the response variable is balanced. . ### convert the response variable into binary 0 means no and 1 means yes

```
Bank$y=ifelse(Bank$y == "yes",1,0)
table(Bank$y)
```

```
##
##    0    1
## 2279 2242
```

Now, the response variable is balanced.

convert the non numeric vairables into factor

```
Bank$job=as.factor(Bank$job)
Bank$marital=as.factor(Bank$marital)
Bank$education=as.factor(Bank$education)
Bank$default=as.factor(Bank$default)
Bank$housing=as.factor(Bank$housing)
Bank$loan=as.factor(Bank$loan)
Bank$contact=as.factor(Bank$contact)
Bank$month=as.factor(Bank$month)
Bank$poutcome=as.factor(Bank$poutcome)
Bank$y=as.factor(Bank$y)
```

summary of the original data

```
summary(Bank)
```



```
##          age                job                marital                education                default
## Min.      :19.00    management :1026    divorced: 591    primary   : 662    no :4432
## 1st Qu.:33.00    blue-collar: 815    married :2647    secondary:2236    yes:  89
## Median :40.00    technician : 718    single  :1283    tertiary  :1443
## Mean   :41.91    admin.      : 465                                unknown   : 180
## 3rd Qu.:50.00    services    : 391
## Max.   :87.00    retired     : 378
##                (Other)    : 728
##          balance    housing    loan                contact                day
## Min.      :-2082    no :2241    no :3956    cellular :3195    Min.      : 1.00
## 1st Qu.:  101    yes:2280    yes: 565    telephone: 371    1st Qu.:  9.00
## Median   :  569                                unknown   : 955    Median :16.00
## Mean     : 1528                                Mean     :15.73
## 3rd Qu.: 1811                                3rd Qu.:21.00
## Max.     :71188                                Max.     :31.00
##
##          month                duration                campaign                pdays
## may       :1117    Min.      :  4    Min.      : 1.000    Min.      : -1.00
## aug       : 676    1st Qu.: 151    1st Qu.: 1.000    1st Qu.: -1.00
## jul       : 627    Median : 273    Median : 2.000    Median : -1.00
## jun       : 487    Mean   : 396    Mean   : 2.572    Mean   : 52.79
## apr       : 384    3rd Qu.: 543    3rd Qu.: 3.000    3rd Qu.: 56.00
## nov       : 371    Max.   :2769    Max.   :32.000    Max.   :871.00
## (Other): 859
##          previous                poutcome                y
## Min.      : 0.0000    failure: 534    0:2279
## 1st Qu.: 0.0000    other   : 251    1:2242
## Median   : 0.0000    success: 362
## Mean     : 0.7554    unknown:3374
## 3rd Qu.: 1.0000
## Max.     :20.0000
##
```

First, let start fitting the full model.

Fit the full model with all the predictors without cross validation, splitting the data and feature selection

```
fullBank=glm(y~.,data=Bank,family="binomial")
summary(fullBank)
```

```
##
## Call:
## glm(formula = y ~ ., family = "binomial", data = Bank)
```

```
##
## Deviance Residuals:
##      Min        1Q      Median        3Q        Max
## -4.2864   -0.5810   -0.1039    0.5901    2.2042
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -1.212e+00  4.548e-01  -2.664 0.007712 **
## age           6.587e-03  5.287e-03   1.246 0.212835
## jobblue-collar -7.218e-01  1.721e-01  -4.193 2.75e-05 ***
## jobentrepreneur -2.708e-01  2.744e-01  -0.987 0.323705
## jobhousemaid   -4.626e-01  3.107e-01  -1.489 0.136525
## jobmanagement -3.298e-01  1.797e-01  -1.835 0.066447 .
## jobretired     1.488e-01  2.259e-01   0.658 0.510236
## jobself-employed -5.089e-01  2.700e-01  -1.885 0.059425 .
## jobservices    -7.598e-01  2.038e-01  -3.728 0.000193 ***
## jobstudent     5.860e-01  2.952e-01   1.985 0.047116 *
## jobtechnician  -5.424e-01  1.678e-01  -3.233 0.001226 **
## jobunemployed  -8.178e-01  2.973e-01  -2.751 0.005941 **
## jobunknown     -1.041e-01  4.514e-01  -0.231 0.817570
## maritalmarried -2.071e-01  1.337e-01  -1.549 0.121328
## maritalsingle  6.200e-02  1.583e-01   0.392 0.695211
## educationsecondary 2.306e-01  1.432e-01   1.610 0.107475
## educationtertiary 3.541e-01  1.672e-01   2.118 0.034139 *
## educationunknown -4.775e-01  2.602e-01  -1.835 0.066435 .
## defaultyes     2.775e-01  3.039e-01   0.913 0.361287
## balance       -9.387e-06  1.464e-05  -0.641 0.521400
## housingyes     -2.609e-01  9.849e-02  -2.649 0.008081 **
## loanyes        -1.005e+00  1.420e-01  -7.078 1.46e-12 ***
## contacttelephone -3.189e-01  1.689e-01  -1.888 0.059089 .
## contactunknown -1.171e+00  1.436e-01  -8.156 3.47e-16 ***
## day           1.246e-02  5.886e-03   2.117 0.034274 *
## monthaug       -5.373e-01  1.742e-01  -3.084 0.002044 **
## monthdec       -2.358e-01  5.969e-01  -0.395 0.692865
## monthfeb       3.898e-01  2.088e-01   1.867 0.061850 .
## monthjan       -1.386e+00  2.797e-01  -4.955 7.22e-07 ***
## monthjul       -9.113e-01  1.847e-01  -4.935 8.01e-07 ***
## monthjun       2.458e-01  2.155e-01   1.140 0.254095
## monthmar       1.611e+00  3.109e-01   5.183 2.19e-07 ***
## monthmay       -9.546e-01  1.726e-01  -5.530 3.21e-08 ***
## monthnov       -9.282e-01  1.986e-01  -4.675 2.94e-06 ***
## monthoct       1.889e+00  2.856e-01   6.615 3.72e-11 ***
## monthsep       7.571e-01  3.741e-01   2.024 0.043019 *
## duration       5.927e-03  2.053e-04  28.875 < 2e-16 ***
## campaign      -1.394e-01  2.262e-02  -6.162 7.20e-10 ***
## pdays         1.625e-03  6.796e-04   2.391 0.016798 *
## previous      -7.190e-02  3.768e-02  -1.908 0.056370 .
```

```
## poutcomeother      8.455e-01  2.065e-01   4.095 4.21e-05 ***
## poutcomesuccess    3.245e+00  3.185e-01  10.187 < 2e-16 ***
## poutcomeunknown    -2.179e-01  2.457e-01  -0.887 0.375051
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6267.1  on 4520  degrees of freedom
## Residual deviance: 3593.2  on 4478  degrees of freedom
## AIC: 3679.2
##
## Number of Fisher Scoring iterations: 6
```

Using the full model, many predictor variables are significant to predict whether or not customer will subscribe the term deposit. This might lead to overfitting since we are using all the features to fit the model.

Just to make sure all the variables are really significant for prediction, I will next use feature selection to identify the most important predictors that contribute to a given outcome variable).

By selecting only the most relevant predictors, we can improve the accuracy and interpretability of predictive models.

In this analysis, I will use forward and backward elimination for feature selection

Backward elimination

```
library(caret)
backwardsModel=step(
  object=fullBank,
  direction = "backward",
  scope=y~.,
  trace=0
)
selectedFeatures=names(coef(backwardsModel))[-1]
print(selectedFeatures)
```

```
## [1] "jobblue-collar"      "jobentrepreneur"    "jobhousemaid"
## [4] "jobmanagement"      "jobretired"         "jobself-employed"
## [7] "jobservices"        "jobstudent"         "jobtechnician"
## [10] "jobunemployed"      "jobunknown"         "maritalmarried"
## [13] "maritalsingle"      "educationsecondary" "educationtertiary"
## [16] "educationunknown"   "housingyes"         "loanyes"
## [19] "contacttelephone"   "contactunknown"     "day"
## [22] "monthaug"           "monthdec"           "monthfeb"
## [25] "monthjan"           "monthjul"           "monthjun"
## [28] "monthmar"           "monthmay"           "monthnov"
## [31] "monthoct"           "monthsep"           "duration"
## [34] "campaign"           "pdays"             "previous"
## [37] "poutcomeother"      "pcomesuccess"       "poutcomeunknown"
```

From the features selection using backward elimination, I found out that variables “job”, “marital”, “education”, “housing”, “loan”, “contact”, “day”, “month”, “duration”, “campaign”, “pdays”, “previous”, “poutcome” are selected features.

forward selection

```
forwardModel=step(
  #fullBank is a original model fit
  object=fullBank,
  direction = "forward",
  scope=y~.,
  trace=0
)
selectedFeatures=names(coef(backwardsModel))[-1]
print(selectedFeatures)
```

```
## [1] "jobblue-collar"      "jobentrepreneur"    "jobhousemaid"
## [4] "jobmanagement"      "jobretired"         "jobself-employed"
## [7] "jobservices"        "jobstudent"         "jobtechnician"
## [10] "jobunemployed"      "jobunknown"         "maritalmarried"
## [13] "maritalsingle"      "educationsecondary" "educationtertiary"
## [16] "educationunknown"   "housingyes"         "loanyes"
## [19] "contacttelephone"   "contactunknown"     "day"
## [22] "monthaug"           "monthdec"           "monthfeb"
## [25] "monthjan"           "monthjul"           "monthjun"
## [28] "monthmar"           "monthmay"           "monthnov"
## [31] "monthoct"           "monthsep"           "duration"
## [34] "campaign"           "pdays"             "previous"
## [37] "poutcomeother"      "pcomesuccess"       "poutcomeunknown"
```

- Backward elimination and forward selection methods choose 13 predictor variables out of 16 from the data.
- From the features selection using forward method, I found out that variables “job”, “marital”, “education”, “housing”, “loan”, “contact”, “day”, “month”, “duration”, “campaign”, “pdays”, “previous”, “poutcome” are selected features.

create a new data from the selected features(it comes from forward selection)

```
selectedFeaturesCol=c("job","marital","education","housing","loan","contact","day","month","duration","campaign","pdays","previous","poutcome")
selected_features <- c(selectedFeaturesCol, "y")
print(selected_features)
```

```
## [1] "job"      "marital"  "education" "housing"  "loan"     "contact"
## [7] "day"      "month"    "duration"  "campaign" "pdays"   "previous"
## [13] "poutcome" "y"
```

```
newBank <- Bank%>%
  dplyr::select(one_of(selected_features))
```

Creating training and testing data with selected features

```
train = sample(dim(newBank)[1], dim(newBank)[1]*0.8)
test=-train
newBank.test=newBank[test,]
newBank.train=newBank[train,]
```

Fit the logistic regression model with selected features using cross-validation.

```
library(caret)
fitControl2 <- trainControl(method = "cv", number = 10)
#fit the modle
newBankFit <- train(y ~ ., data = newBank.train , method = "glm", trControl = fitControl2,family=binomial)
#predicitons
predictions2=predict(newBankFit,newdata = newBank.test,type="prob")
binaryPreds2=ifelse(predictions2[,2]>0.5,1,0)
binaryPredsfactor2=factor(binaryPreds2,levels=c(0,1),ordered=TRUE)
```

confusion matrix with selected features from forward selection

```
confusion2=confusionMatrix(binaryPredsfactor2,newBank.test$y)
confusion2
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0 377  78
##              1  78 372
##
##              Accuracy : 0.8276
##              95% CI : (0.8014, 0.8517)
##      No Information Rate : 0.5028
##      P-Value [Acc > NIR] : <2e-16
##
##              Kappa : 0.6552
##
##  McNemar's Test P-Value : 1
##
##              Sensitivity : 0.8286
##              Specificity : 0.8267
##              Pos Pred Value : 0.8286
##              Neg Pred Value : 0.8267
##              Prevalence : 0.5028
##              Detection Rate : 0.4166
##      Detection Prevalence : 0.5028
##              Balanced Accuracy : 0.8276
##
##              'Positive' Class : 0
##
```

ROC curve with selected features from forward selection


note: roc() function expects predicted class probabilities, not class labels

```
rocobj1 <- roc(newBank.test$y, binaryPredsfactor2)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(rocobj1, print.auc=TRUE, print.auc.x=0.5, print.auc.y=0.2)
```



```
aucScore1=auc(rocobj1)  
print(aucScore1)
```

```
## Area under the curve: 0.8276
```

AUC=0.828. It is greater than 0.5. It means the model perform better than random guessing.

calculate precision and accuracy

```
truePositive2=confusion2$table[2,2]
falsePositive2=confusion2$table[1,2]
precision2=truePositive2/(truePositive2+falsePositive2)
accuracy2=confusion2$overall["Accuracy"]
cat("Precision measures how often the model correctly predicts that customers will su
bscribe the term deposit.")
```

```
## Precision measures how often the model correctly predicts that customers will subs
cribe the term deposit.
```

```
cat("\n")
```

```
cat("\n")
```

```
cat("The precision score using cross validation with the selected feature is"
,round(precision2*100,2), "%.")
```

```
## The precision score using cross validation with the selected feature is 82.67 %.
```

```
cat("\n")
```

```
cat("\n")
```

```
cat("Accuracy measures how often the model correctly predicts, regardless of it is ab
out predicting no subscribe or subscribe the term deposit.")
```

```
## Accuracy measures how often the model correctly predicts, regardless of it is abou
t predicting no subscribe or subscribe the term deposit.
```

```
cat("\n")
```



```
cat("The accuracy score using cross validation with the selected feature is"  
,round(accuracy2*100,2), "%")
```

```
## The accuracy score using cross validation with the selected feature is 82.76 %
```

Next fit the model using the random Forest

Advantages: More stable and robust than decision trees, can handle both continuous and categorical features, handles irrelevant features well, can capture non-linear relationships between the input features and the output, can handle large datasets.

Disadvantages: Can be slow and memory-intensive, may not perform well on imbalanced datasets, may not work well with high-dimensional data.

fit the model using random Forest on selected features with cross validation

```
random.fit=train(y~.,data=newBank.train,method="rf",trControl=fitControl2)  
predictions3=predict(random.fit,newdata = newBank.test,type="prob")  
binaryPreds3=ifelse(predictions3[,2]>0.5,1,0)  
binaryPredsfactor3=factor(binaryPreds3,levels=c(0,1),ordered=TRUE)
```

confusion matrix with selected features from forward selection

```
confusion3=confusionMatrix(binaryPredsfactor3,newBank.test$y)  
confusion3
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 431    3
##           1   24 447
##
##           Accuracy : 0.9702
##           95% CI : (0.9569, 0.9802)
##           No Information Rate : 0.5028
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9403
##
## Mcnemar's Test P-Value : 0.0001186
##
##           Sensitivity : 0.9473
##           Specificity : 0.9933
##           Pos Pred Value : 0.9931
##           Neg Pred Value : 0.9490
##           Prevalence : 0.5028
##           Detection Rate : 0.4762
##           Detection Prevalence : 0.4796
##           Balanced Accuracy : 0.9703
##
##           'Positive' Class : 0
##
```

ROC curve with selected features from forward selection


note: roc() function expects predicted class probabilities, not class labels

```
rocobj2 <- roc(newBank.test$y, binaryPredsfactor3)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(rocobj2, print.auc=TRUE, print.auc.x=0.5, print.auc.y=0.2)
```



```
aucScore2=auc(rocobj2)
print(aucScore2)
```

```
## Area under the curve: 0.9703
```

AUC=0.967. It is greater than 0.5. It means the model perform better than random guessing. ### calculate precision and accuracy

```
truePositive3=confusion3$table[2,2]
falsePositive3=confusion3$table[1,2]
precision3=truePositive3/(truePositive3+falsePositive3)
accuracy3=confusion3$overall["Accuracy"]
cat("Precision measures how often the model correctly predicts that customers will su
bscribe the term deposit.")
```

```
## Precision measures how often the model correctly predicts that customers will subs
cribe the term deposit.
```

```
cat("\n")
```

```
cat("The precision score using cross validation with the selected feature is",  
round(precision3*100,2), "%.")
```

```
## The precision score using cross validation with the selected feature is 99.33 %.
```

```
cat("\n")
```

```
cat("\n")
```

```
cat("Accuracy measures how often the model correctly predicts, regardless of it is ab  
out predicting no subscribe or subscribe the term deposit.")
```

```
## Accuracy measures how often the model correctly predicts, regardless of it is abou  
t predicting no subscribe or subscribe the term deposit.
```

```
cat("\n")
```

```
cat("The accuracy score using cross validation with the selected feature is",  
round(accuracy3*100,2), "%")
```

```
## The accuracy score using cross validation with the selected feature is 97.02 %
```

Random Forest significantly increase the accuracy and auc in the new bank data. However, it is computationally expensive.

Fit the model using Decision trees. It can handle both classification and regression tasks.

some advantages and disadvantages of using Decision trees.

Advantages: It can handle nonlinear relationships and mixed feature types.

Disadvantages: It is prone to overfitting. It is sensitivity to small variations.

Decision trees

```
treeModel=train(y ~ ., data = newBank.train , method = "rpart", trControl = fitContro
l2)
predictions4=predict(treeModel,newdata = newBank.test)
```

confusion matrix with selected features from forward selection

```
confusion4=confusionMatrix(data=predictions4,newBank.test$y)
confusion4
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0 329   72
##              1 126  378
##
##              Accuracy : 0.7812
##              95% CI : (0.7528, 0.8078)
##              No Information Rate : 0.5028
##              P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.5627
##
## Mcnemar's Test P-Value : 0.0001655
##
##              Sensitivity : 0.7231
##              Specificity : 0.8400
##              Pos Pred Value : 0.8204
##              Neg Pred Value : 0.7500
##              Prevalence : 0.5028
##              Detection Rate : 0.3635
##              Detection Prevalence : 0.4431
##              Balanced Accuracy : 0.7815
##
##              'Positive' Class : 0
##
```

ROC curve with selected features from forward selection


note: roc() function expects predicted class probabilities, not class labels

```
response=as.numeric(newBank.test$y)-1  
predictor=as.numeric(predictions4)-1  
rocobj3 <- roc(response, predictor)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(rocobj3,print.auc=TRUE,print.auc.x=0.5,print.auc.y=0.2)
```



```
aucScore3=auc(rocobj3)  
print(aucScore3)
```

```
## Area under the curve: 0.7815
```

calculate precision and accuracy

```
truePositive4=confusion4$table[2,2]
falsePositive4=confusion4$table[1,2]
precision4=truePositive4/(truePositive4+falsePositive4)
accuracy4=confusion4$overall["Accuracy"]
cat("Precision measures how often the model correctly predicts that customers will su
bscribe the term deposit.")
```

```
## Precision measures how often the model correctly predicts that customers will subs
cribe the term deposit.
```

```
cat("\n")
```

```
cat("The precision score using cross validation with the selected feature is"
,round(precision4*100,2), "%.")
```

```
## The precision score using cross validation with the selected feature is 84 %.
```

```
cat("\n")
```

```
cat("\n")
```

```
cat("Accuracy measures how often the model correctly predicts, regardless of it is ab
out predicting no subscribe or subscribe the term deposit.")
```

```
## Accuracy measures how often the model correctly predicts, regardless of it is abou
t predicting no subscribe or subscribe the term deposit.
```

```
cat("\n")
```

```
cat("The accuracy score using cross validation with the selected feature is"
,round(accuracy4*100,2), "%")
```

```
## The accuracy score using cross validation with the selected feature is 78.12 %
```

Next, fit the model using Naive Bayes.

The only problem of using Naive Bayes is that it is assume that features are conditionally independent.

```
naiveBayesFit=train(y ~ ., data = newBank.train , method = "naive_bayes", trControl =
fitControl2)
predictions5=predict(naiveBayesFit,newdata = newBank.test)
```

confusion matrix with selected features from forward selection

```
confusion5=confusionMatrix(data=predictions5,newBank.test$y)
confusion5
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0 392 220
##              1   63 230
##
##              Accuracy : 0.6873
##              95% CI : (0.656, 0.7174)
##      No Information Rate : 0.5028
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.3734
##
##  McNemar's Test P-Value : < 2.2e-16
##
##              Sensitivity : 0.8615
##              Specificity : 0.5111
##      Pos Pred Value : 0.6405
##      Neg Pred Value : 0.7850
##      Prevalence : 0.5028
##      Detection Rate : 0.4331
##      Detection Prevalence : 0.6762
##      Balanced Accuracy : 0.6863
##
##              'Positive' Class : 0
##
```

ROC curve with selected features from forward selection


note: roc() function expects predicted class probabilities, not class labels


```
response2=as.numeric(newBank.test$y)-1  
predictor2=as.numeric(predictions5)-1  
rocobj4 <- roc(response2, predictor2)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(rocobj4,print.auc=TRUE,print.auc.x=0.5,print.auc.y=0.2)
```



```
aucScore4=auc(rocobj4)  
print(aucScore4)
```

```
## Area under the curve: 0.6863
```

calculate precision and accuracy

```
truePositive5=confusion5$table[2,2]
falsePositive5=confusion5$table[1,2]
precision5=truePositive5/(truePositive5+falsePositive5)
accuracy5=confusion5$overall["Accuracy"]
cat("Precision measures how often the model correctly predicts that customers will su
bscribe the term deposit.")
```

```
## Precision measures how often the model correctly predicts that customers will subs
cribe the term deposit.
```

```
cat("\n")
```

```
cat("The precision score using cross validation with the selected feature is"
,round(precision5*100,2), "%.")
```

```
## The precision score using cross validation with the selected feature is 51.11 %.
```

```
cat("\n")
```

```
cat("\n")
```

```
cat("Accuracy measures how often the model correctly predicts, regardless of it is ab
out predicting no subscribe or subscribe the term deposit.")
```

```
## Accuracy measures how often the model correctly predicts, regardless of it is abou
t predicting no subscribe or subscribe the term deposit.
```

```
cat("\n")
```

```
cat("The accuracy score using cross validation with the selected feature is"
,round(accuracy5*100,2), "%")
```

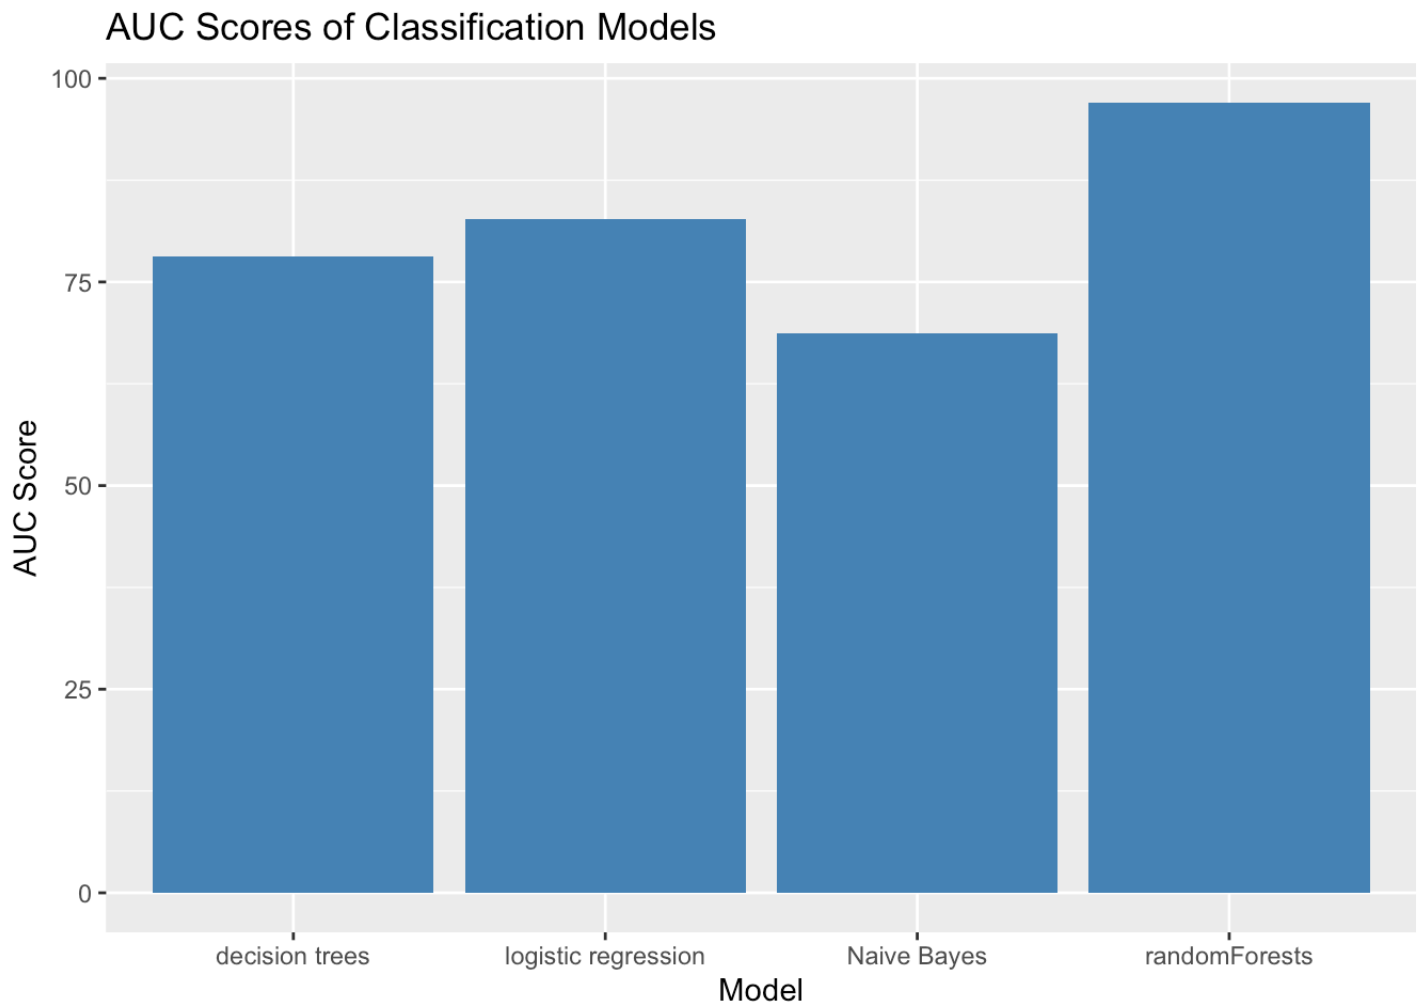
```
## The accuracy score using cross validation with the selected feature is 68.73 %
```

Create a data frame containing the model names and corresponding AUC, accuracy, and precision scores.

```
modelScores=data.frame(Model=c("logistic regression", "randomForests", "decision trees", "Naive Bayes"),
                        Precision=c(round(precision2*100,2),round(precision3*100,2),round(precision4*100,2),round(precision5*100,2)),
                        Accuracy=c(round(accuracy2 *100,2),round(accuracy3 *100,2),round(accuracy4*100,2),round(accuracy5*100,2)),
                        AUC=c(round(aucScore1 *100,2),round(aucScore2 * 100,2),round(aucScore3 * 100,2),round(aucScore4*100,2)))
```


Create a bar graph showing the AUC scores of logistic regression, random Forest, and decision trees

```
ggplot(modelScores,aes(x=Model,y=AUC))+
  geom_bar(stat="identity",fill="steelblue")+
  labs(title = "AUC Scores of Classification Models",
       x= "Model",y="AUC Score")
```



Create a bar graph showing the Precision scores of logistic regression, random Forest, and decision trees

```
ggplot(modelScores,aes(x=Model,y=Precision))+  
  geom_bar(stat="identity",fill="steelblue")+  
  labs(title = "Precision Scores of Classification Models",  
        x= "Model",y="Precision Score")
```



Create a bar graph showing the Accuracy scores of logistic regression, random Forest, and decision trees

```
ggplot(modelScores,aes(x=Model,y=Accuracy))+  
  geom_bar(stat="identity",fill="steelblue")+  
  labs(title = "Accuracy Scores of Classification Models",  
        x= "Model",y="Accuracy Score")
```

